

COMP 6660 Fall 2020 Assignment 1C

Jordan Sosnowski

jjis0029@auburn.edu

Table of Contents

Green 3:	3
Performance	5
BC1	5
BC2	6
C1	6
C2	7
Green 3.1: 1B Evolutionary Algorithm vs 1C Evolutionary Algorithm	8
1BC	8
2BC	8
Green 3.2: Validity Forced Analysis	9
1BC	9
2BC	10
C1	11
C2	11
Green 3.3: Exploring Penalty Coefficients	12
10 vs 5	12
50 v 10.....	13
50 v 5.....	14
Yellow 1: Self-adaptive Mutation Rate	15
Red 2: Self-adaptive Penalty Coefficient.....	16

Green 3:

For this assignment series, we were asked to implement constraint-satisfaction for our evolutionary algorithm. Specifically, a penalty function for constraints. Instead of immediately invalidating solutions and giving them a fitness of 0 a penalty function would be applied to them. My penalty function was as follows. First, you iterate through a dictionary of black cells that are not satisfied. You then get the difference in what the value of bulbs around the cell should be vs what was and take the absolute value of that. This is added to a penalty sum. After checking unsatisfied black cells, I go on to check the amount of intersecting light rays there were. Each number of intersecting rays are added up to the penalty sum as well. Finally, this penalty sum is multiplied by the penalty coefficient and subtracted from the original fitness.

```
def penalty_function(self, fitness, black_cells, intersections, penalty_coef):
    """ Penalizes current individual based on the amount of intersecting rays
    and black cells that are not constrained.

    f' = f - p*(black_cells + intersections)

    where f is the original fitness
    where f' is the new fitness
    where p is the penalty coefficient
    """

    penalty = 0
    for expected, actual_list in black_cells.items():
        for actual in actual_list:
            penalty += abs(expected - actual)

    for section in intersections.values():
        for intersection in section:
            penalty += intersection

    fitness -= penalty*penalty_coef
    return fitness
```

Using a penalty function was also easy to integrate as I simply inserted it into the fitness function but with a conditional statement based on a user-defined parameter.

```
def calculate_fitness(self, board, penalty_coef):
    """ Calculates fitness of current board.

    board - RxC numpy array representing the game board, filled with bulbs,
    black cells, and light
    """

    fitness = lightup.calculate_completion(board)
    black_cells = lightup.check_black_cells(board, self.ignore_black_cells)
    intersections = lightup.check_intersections(board)
    fitness_prime = fitness

    if self.constraint_algorithm in CON_PENALTY:
        if penalty_coef is None:
            penalty_coef = self.penalty_coefficient
        fitness_prime = self.penalty_function(
            fitness, black_cells, intersections, penalty_coef)

    penalty = fitness - fitness_prime
    return fitness_prime, penalty
```

We are additionally asked to implement a new survival strategy. Thus far we have been using a “plus” survival strategy or a $(\mu+\lambda)$ -EA. Now we were asked to implement a comma or generational EA $[(\mu,\lambda)$ -EA]. For a generational EA for every generation, the existing generation is discarded and only the children live on. Due to the modularity of my code, this was rather easy to implement. After the children are created, I call the below function to determine which individuals will be passed to survival selection.

```
def survival_strategy_selection(self, old_generation, children):
    """ Decide on next generation for survival function.

    old_generation - list of individuals that make up the older generation
    children - list of individuals that make up children for the new generation
    """

    if self.survival_strategy_algorithm in PLUS:
        #  $(\mu + \lambda)$ -EA
        return old_generation + children

    if self.survival_strategy_algorithm in COMMA:
        #  $(\mu, \lambda)$ -EA
        return children
```

In addition to the survival strategy, we were asked to implement a new parent selection algorithm and two new survival selection algorithms. First, I implemented a uniform random parent selection algorithm which was trivial.

```
def uniform_random_parent(self, individuals):
    """ Selects parents randomly with a uniform distribution

    individuals - list of Individual objects
    """

    mating_pool = random.choices(individuals, k=self.children*2)
    return mating_pool
```

Following this, I implemented the uniform random for survival selection as well. Since for survival, we do not want to sample with the replacement I used `random.sample` as `random.choices` use replacement.

```
def uniform_random_survival(self, individuals):
    """ Selects parents randomly without replacement with a uniform distribution

    individuals - list of Individual objects
    """

    mating_pool = random.sample(individuals, k=self.parents)
    return mating_pool
```

Last was FPS for survival selection. For this I had to use NumPy's `random.choice` as Python's standard library's `sample` does not allow weight to be passed to it.

```
def fitness_proportional_selection_survival(self, individuals):
    """ Selects parents based on fitness probability without replacement.

    individuals - list of Individual objects
    """

    fitness_sum = sum(individuals)
    if fitness_sum == 0:
        return random.choices(individuals, k=self.children*2)
```

```
        probs = [individual.fitness /  
                  fitness_sum for individual in individuals]  
  
        return list(np.random.choice(individuals, p=probs, size=self.children*2, replace=False)  
    )
```

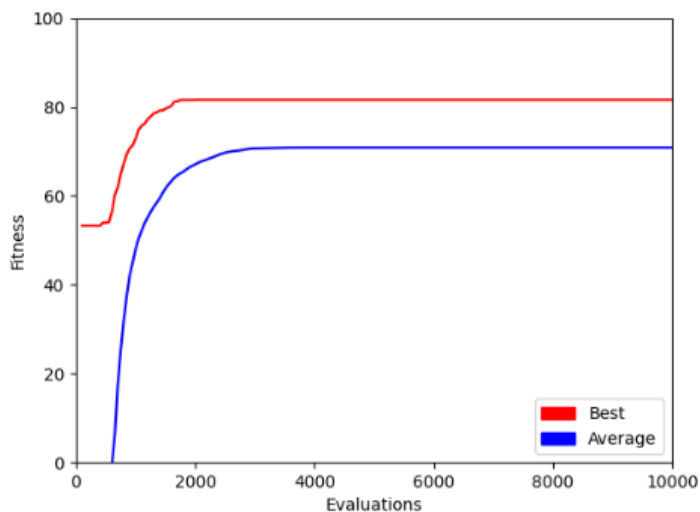
Performance

For the problems for this assignment, I used the same EA parameters for each one. They are listed below:

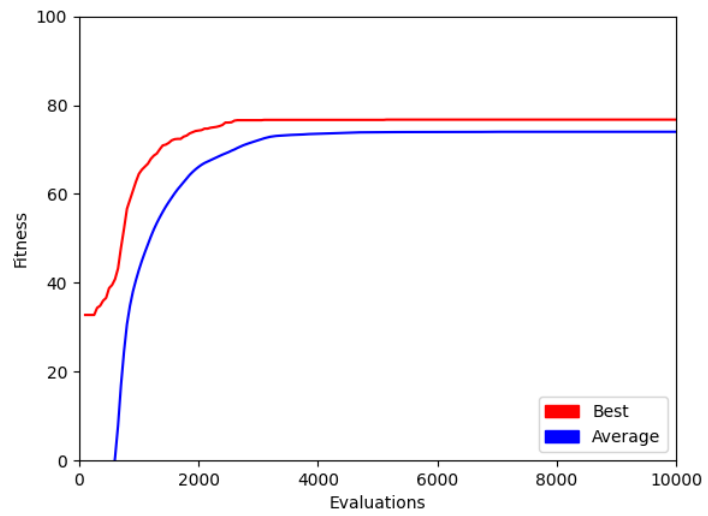
```
{  
    "children": 50,  
    "parents": 100,  
    "parent_alg": "sus",  
    "child_alg": "one point crossover",  
    "survival_alg": "truncation",  
    "termination_alg": "num of evals",  
    "num_of_runs": 30,  
    "initialization_alg": "vanilla",  
    "survival_strategy_alg": "plus",  
    "fitness_evals": 10000,  
    "mutation_rate": 0.40,  
    "penalty_coefficient": 10,  
    "constraint": "penalty",  
    "ignore_black_cells": false,  
    "log_1b": true  
}
```

As BC# are easier problems the graphs are what I expected. Since C# are more difficult their poor graphs are also not surprising. I think I could achieve higher performance on those but since that is not the focus on these assignments I did not put to much effort into those portions.

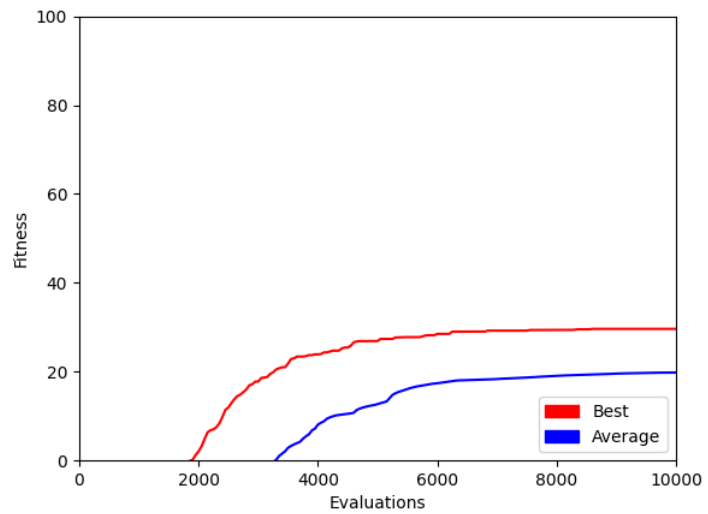
BC1



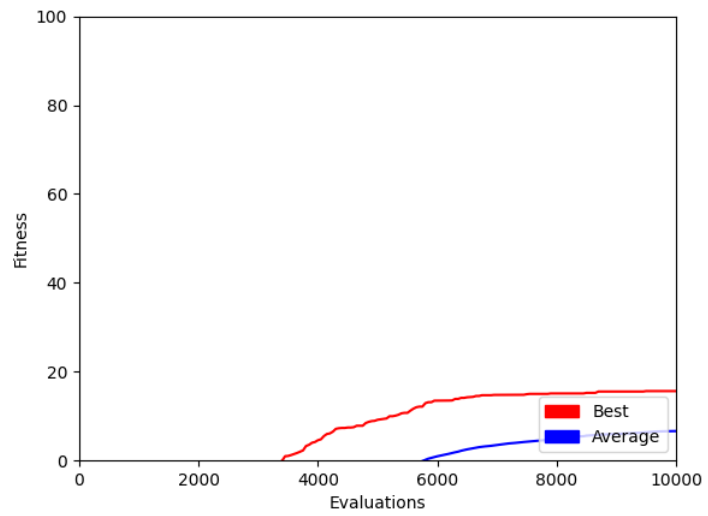
BC2



C1



C2



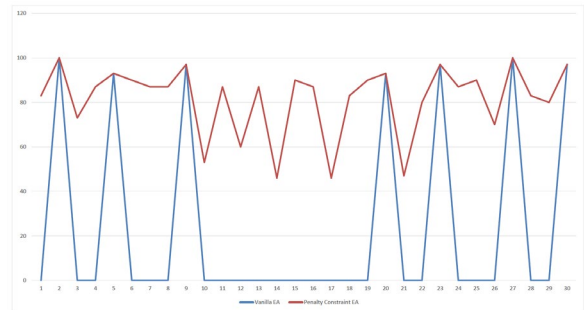
Green 3.1: 1B Evolutionary Algorithm vs 1C Evolutionary Algorithm

1BC

Here we have a comparison against 1B's EA and 1C's EA. 1B is using a "vanilla" EA and 1C is using a penalty constraint EA. As you can see 1C has higher fitnesses where 1C mostly has fitnesses of 0. Unfortunately, this means that most of the solutions in 1C are technically invalid but if this EA could run longer in theory, we would see these penalized solutions go away and be replaced with "valid" ones.

Vanilla EA	Penalty Constraint EA	F-Test Two-Sample for Variances	
0	83		
100	100		
0	73		
0	87		
93	93		
0	90		
0	87		
0	87		
97	97		
0	53		
0	87		
0	60		
0	46		
0	90		
0	87		
0	46		
0	83		
0	90		
93	93		
0	47		
0	80		
97	97		
0	87		
0	90		
0	70		
100	100		
0	83		
0	80		
97	97		

Variable 1	Variable 2
Mean	22.56666667 81.66666667
Variance	1732.667816 254.1609195
Observations	30 30
df	29 29
F	6.817207851
P(F<=f) one-tail	7.98539E-07
F Critical one-tail	1.860811435
F > F-Critical	
assume equal	
t-Test: Two-Sample Assuming Equal Variances	
Variable 1	Variable 2
Mean	22.56666667 81.66666667
Variance	1732.667816 254.1609195
Pooled Variance	993.4143678
Hypothesized Mean Difference	0
df	58
t Stat	-7.262194763
P(T<=t) one-tail	5.31389E-10
t Critical one-tail	1.671552762
P(T<=t) two-tail	1.06278E-09
t Critical two-tail	2.001717484

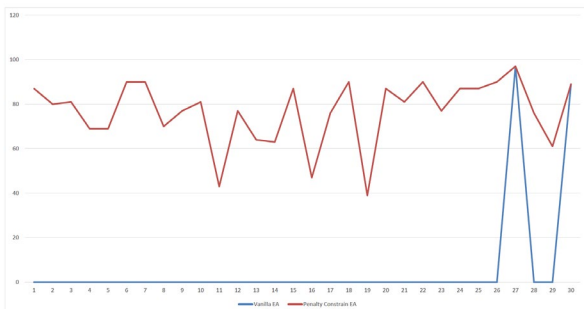


2BC

Here we can see the same thing except with the problem 2BC.

Vanilla EA	Penalty Constrains EA	F-Test Two-Sample for Variances	
0	87		
0	80		
0	81		
0	69		
0	69		
0	90		
0	90		
0	70		
0	77		
0	81		
0	43		
0	77		
0	64		
0	63		
0	87		
0	47		
0	76		
0	90		
0	39		
0	87		
0	81		
0	90		
0	77		
0	87		
0	87		
0	90		
97	97		
0	76		
0	61		
89	89		

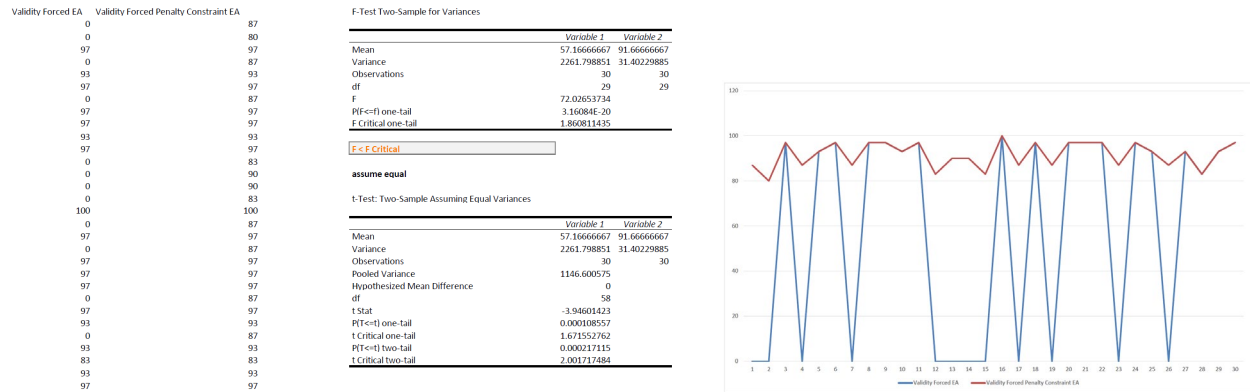
Variable 1	Variable 2
Mean	6.2 76.73333333
Variance	557.8206897 215.3057471
Observations	30 30
df	29 29
F	2.590830468
P(F<=f) one-tail	0.006280087
F Critical one-tail	1.860811435
F > F Critical	
assume equal variance	
t-Test: Two-Sample Assuming Equal Variances	
Variable 1	Variable 2
Mean	6.2 76.73333333
Variance	557.8206897 215.3057471
Pooled Variance	386.5632184
Hypothesized Mean Difference	0
df	58
t Stat	-13.89407932
P(T<=t) one-tail	2.08361E-20
t Critical one-tail	1.671552762
P(T<=t) two-tail	4.16721E-20
t Critical two-tail	2.001717484



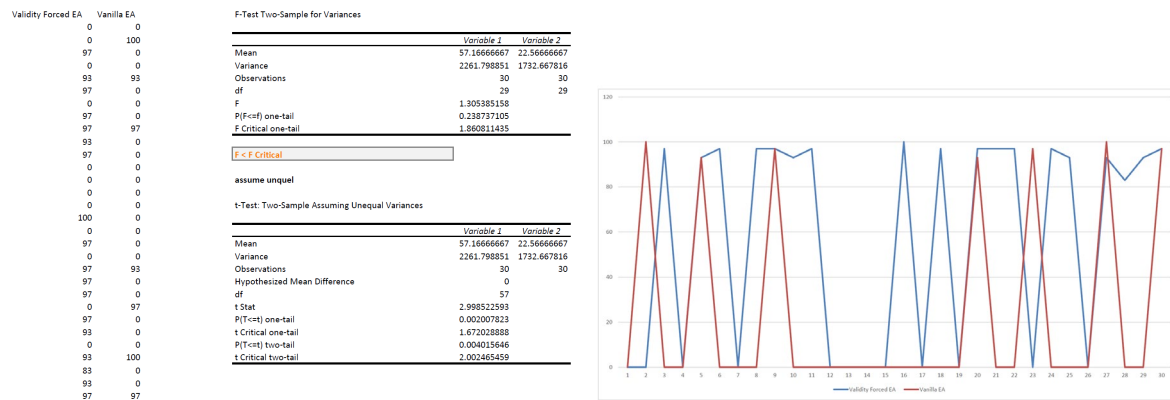
Green 3.2: Validity of Forced Analysis

1BC

For this analysis, we have a Validity Forced Vanilla EA against a Validity Forced Penalty Constraint EA. With the same logic as before we expect to see the penalty constraint to perform better.



For the next analysis, we compare validity forced vs a plain vanilla EA. Since the validity is enforced, we see that the Validity Forced EA is performing slightly better than the plain Vanilla EA which is to be expected.



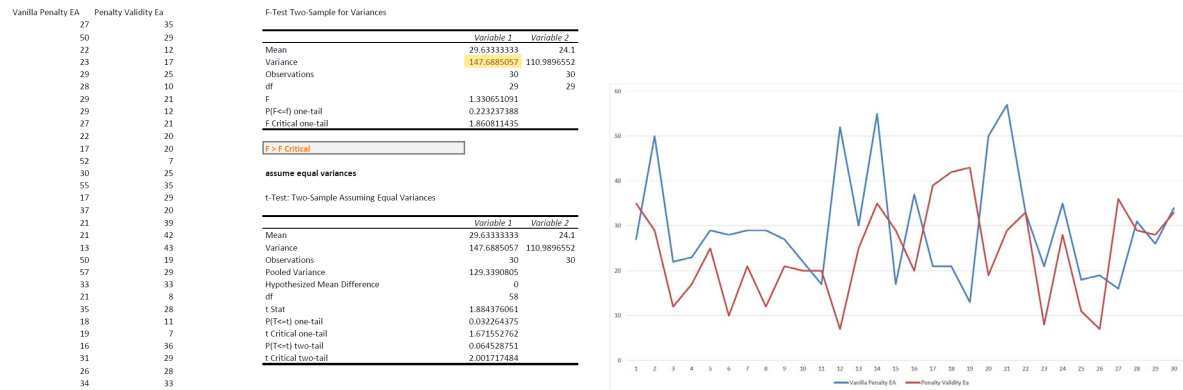
2BC

Here we can see the same thing except with the problem 2BC.

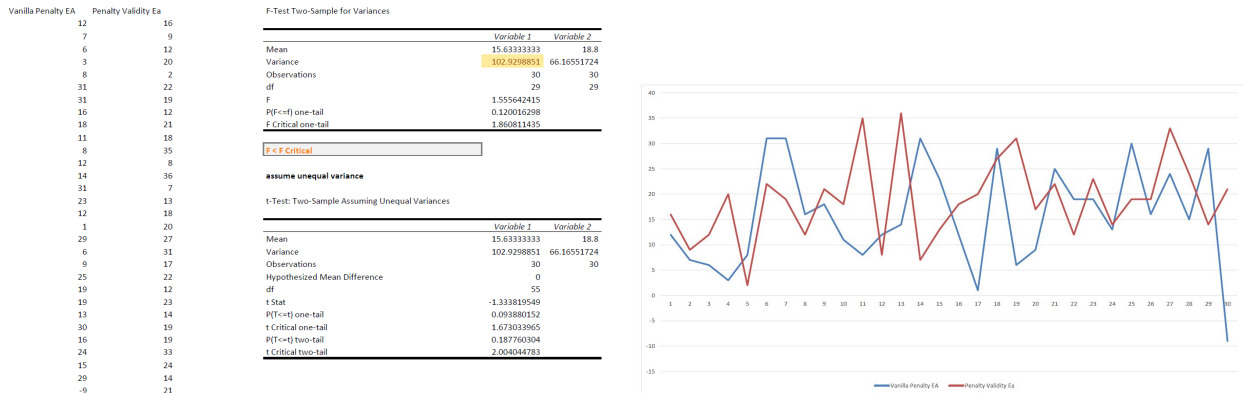


C1

Here we can see the same thing except with the problem C1 which is meant to be harder than the BC problems.



C2



Green 3.3: Exploring Penalty Coefficients

For the last green objective, we were asked to explore penalty coefficients. When I first implemented the penalty coefficient, I picked 10 as that seemed like a good number, and the results, I got were not awful, so I stuck with it. For this objective, I picked a coefficient that was lower and was significantly higher; 5 and 50. My main question after doing this was what is the best way of finding a good penalty coefficient? As if you make it too high it will just kill everything, but if its too low the invalid solutions will look “good”. Additionally, what is the best way of checking that information? In our logs we only check fitness but if we have a low penalty it will look like everything has a “good” fitness.

10 vs 5

We can see that a penalty of 5 outperforms a penalty coefficient of 10. However, as discussed later a lower penalty coefficient can allow seriously “invalid” solutions to flourish. So I think it would most likely be best to be a little higher but with lower performance as the solutions are most likely similar.

C1

10 Coef	5 Coef
27	57
50	57
22	45
23	42
29	53
28	50
29	56
29	54
27	73
22	66
17	47
52	58
30	49
55	54
17	52
37	58
21	64
21	44
13	62
50	49
57	59
33	49
21	57
35	52
18	47
19	56
16	56
31	52
26	56
34	57

F-Test Two-Sample for Variances

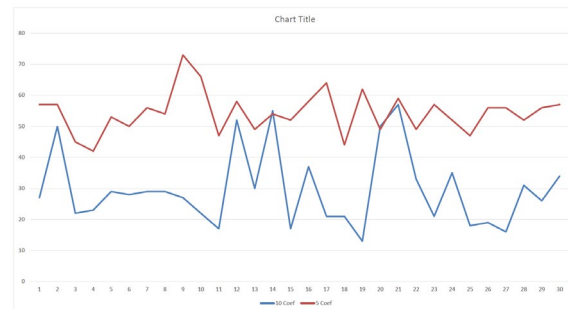
	Variable 1	Variable 2
Mean	29.63333333	54.36666667
Variance	147.6885057	44.86091954
Observations	30	30
df	29	29
F	3.292141741	
P(F<=f) one-tail	0.000986375	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	29.63333333	54.36666667
Variance	147.6885057	44.86091954
Observations	30	30
Pooled Variance	96.27471264	
Hypothesized Mean Difference	0	
df	58	
t Stat	-9.762749917	
P(T<=t) one-tail	3.74972E-14	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	7.49945E-14	
t Critical two-tail	2.001717484	



C2

10 Coef	5 Coef
12	48
7	53
6	54
3	50
8	46
31	58
31	65
16	55
18	66
11	57
8	48
12	53
14	46
31	52
23	53
12	53
1	48
29	48
6	48
9	47
25	53
19	53
19	51
13	53
30	51
16	49
24	46
15	49
29	58
-9	57

F-Test Two-Sample for Variances

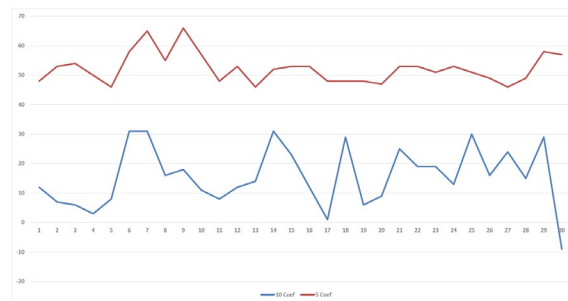
	Variable 1	Variable 2
Mean	15.63333333	52.26666667
Variance	102.9298851	25.5816092
Observations	30	30
df	29	29
F	4.023589145	
P(F<=f) one-tail	0.000170529	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	15.63333333	52.26666667
Variance	102.9298851	25.5816092
Observations	30	30
Pooled Variance	64.25574713	
Hypothesized Mean Difference	0	
df	58	
t Stat	-17.6997071	
P(T<=t) one-tail	2.3494E-25	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	4.69881E-25	
t Critical two-tail	2.001717484	



50 v 10

Here we have a coefficient of 50 running against 10. Logically it makes sense that the coefficient of 50 is seriously underperforming. As one mistake would have a fitness of 50 and two mistakes would have a fitness of 0 (assuming the base fitness was 100).

C1

50 Coef 10 Coef
-227 27
-186 50
-213 22
-224 23
-233 29
-76 28
-219 29
-220 29
-212 27
-173 22
-240 17
-115 52
-167 30
-170 55
-264 17
-221 37
-158 21
-232 21
-212 13
-162 50
-180 57
-125 33
-290 21
-185 35
-270 18
-179 19
-180 16
-258 31
-183 26
-219 34

F-Test Two-Sample for Variances

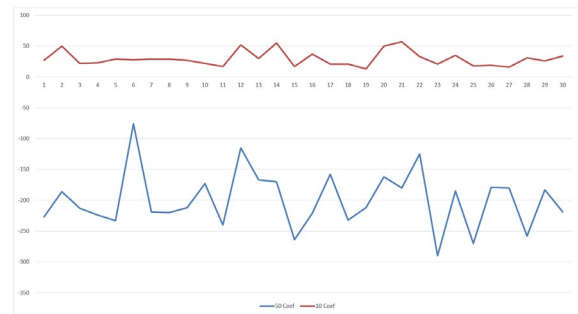
	Variable 1	Variable 2
Mean	-199.7666667	29.63333333
Variance	2193.081609	147.6885057
Observations	30	30
df	29	29
F	14.84937232	
P(F<=f) one-tail	6.91186E-11	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	-199.7666667	29.63333333
Variance	2193.081609	147.6885057
Observations	30	30
Pooled Variance	1170.385057	
Hypothesized Mean Difference	0	
df	58	
t Stat	-25.97016201	
P(T<=t) one-tail	6.24425E-34	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	1.24885E-33	
t Critical two-tail	2.001717484	



C2

50 Coef 10 Coef
-281 12
-68 7
-162 6
-220 3
-217 8
-214 31
-165 31
-220 16
-213 18
-224 11
-212 8
-317 12
-169 14
-173 31
-177 23
-212 12
-275 1
-223 29
-111 6
-227 9
-182 25
-274 19
-227 19
-173 13
-211 30
-271 16
-235 24
-223 15
-166 29
-217 -9

F-Test Two-Sample for Variances

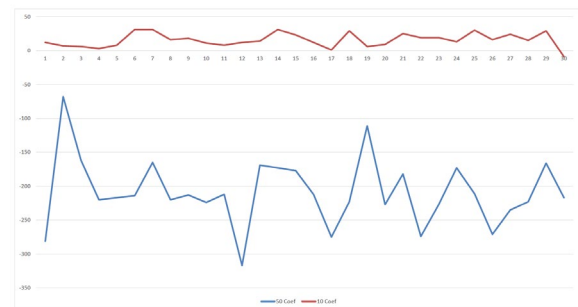
	Variable 1	Variable 2
Mean	-208.6333333	15.63333333
Variance	2519.550575	102.9298851
Observations	30	30
df	29	29
F	24.47831913	
P(F<=f) one-tail	9.71105E-14	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	-208.6333333	15.63333333
Variance	2519.550575	102.9298851
Observations	30	30
Pooled Variance	1311.24023	
Hypothesized Mean Difference	0	
df	58	
t Stat	-23.98662939	
P(T<=t) one-tail	4.25626E-32	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	8.51253E-32	
t Critical two-tail	2.001717484	



50 v 5

As seen earlier 5 performs the “best” so it is no surprise that it outperforms 50.

C1

50 Coef 5 Coef

-227 57
-186 57
-213 45
-224 42
-233 53
-76 50
-219 56
-220 54
-212 73
-173 66
-240 47
-115 58
-167 49
-170 54
-264 52
-221 58
-158 64
-232 44
-212 62
-162 49
-180 59
-125 49
-290 57
-185 52
-270 47
-179 56
-180 56
-258 52
-183 56
-219 57

F-Test Two-Sample for Variances

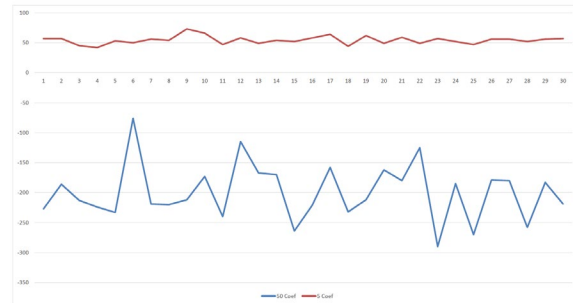
	Variable 1	Variable 2
Mean	-199.7666667	54.36666667
Variance	2193.081609	44.86091954
Observations	30	30
df	29	29
F	48.88623844	
P(F<=f) one-tail	7.3145E-18	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	-199.7666667	54.36666667
Variance	2193.081609	44.86091954
Observations	30	30
Pooled Variance	1118.971264	
Hypothesized Mean Difference	0	
df	58	
t Stat	-29.42373466	
P(T<=t) one-tail	7.39895E-37	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	1.47979E-36	
t Critical two-tail	2.001717484	



C2

50 Coef 5 Coef

-281 48
-68 53
-162 54
-220 50
-217 46
-214 58
-165 65
-220 55
-213 66
-224 57
-212 48
-317 53
-169 46
-173 52
-177 53
-212 53
-275 48
-223 48
-111 48
-227 47
-182 53
-274 53
-227 51
-173 53
-211 51
-271 49
-235 46
-223 49
-166 58
-217 57

F-Test Two-Sample for Variances

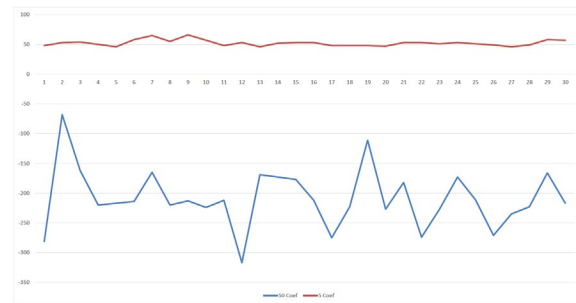
	Variable 1	Variable 2
Mean	-208.6333333	52.26666667
Variance	2519.550575	25.5816092
Observations	30	30
df	29	29
F	98.49069914	
P(F<=f) one-tail	3.73781E-22	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	-208.6333333	52.26666667
Variance	2519.550575	25.5816092
Observations	30	30
Pooled Variance	1272.566092	
Hypothesized Mean Difference	0	
df	58	
t Stat	-28.3256272	
P(T<=t) one-tail	5.83337E-36	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	1.16667E-35	
t Critical two-tail	2.001717484	



Yellow 1: Self-adaptive Mutation Rate

A self-adaptive mutation rate, as far as I understand, is having the mutation rate being defined as a gene that can be passed between parents instead of something that is defined to be static for the entire experiment. As I used a class-based approach for my Individuals all I had to do was add a new instance variable to it. Then during recombination, mutation, and initialization I had to ensure to generate/pick/mutate/use that new gene. I did expect the self-adaptive mutation to do better but they perform very similar.

C1

self-adaptivity mutation	Base	F-Test Two-Sample for Variances
36	35	
22	29	
15	12	
37	17	
25	25	
0	10	
13	21	
7	12	
28	21	
28	20	
36	20	
46	7	
23	25	
33	35	
13	29	
29	20	
7	39	
19	42	
20	43	
50	19	
4	29	
22	33	
8	8	
6	28	
19	11	
33	7	
19	36	
18	29	
18	28	
21	33	

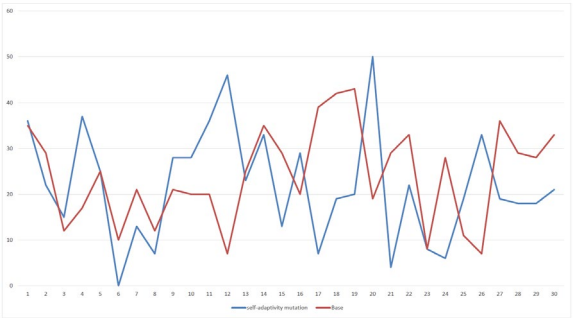
Variable 1	Variable 2
Mean	21.83333333
Variance	149.454023
Observations	30
df	29
F	1.346558134
P(F<=f) one-tail	0.213929127
F Critical one-tail	1.860811435

F < F Critical

assume unequal variance

t-Test: Two-Sample Assuming Unequal Variances

Variable 1	Variable 2
Mean	21.83333333
Variance	149.454023
Observations	30
Hypothesized Mean Difference	0
df	57
t Stat	-0.76929228
P(T<=t) one-tail	0.222448018
t Critical one-tail	1.672028888
P(T<=t) two-tail	0.444896037
t Critical two-tail	2.002465459



C2

self-adaptivity mutation	base	F-Test Two-Sample for Variances
21	16	
23	9	
16	12	
26	20	
35	2	
25	22	
18	19	
27	12	
14	21	
29	18	
24	35	
2	8	
14	36	
15	7	
0	13	
29	18	
11	20	
37	27	
18	31	
15	17	
15	22	
24	12	
21	23	
5	14	
17	19	
13	19	
24	33	
13	24	
43	14	
23	21	

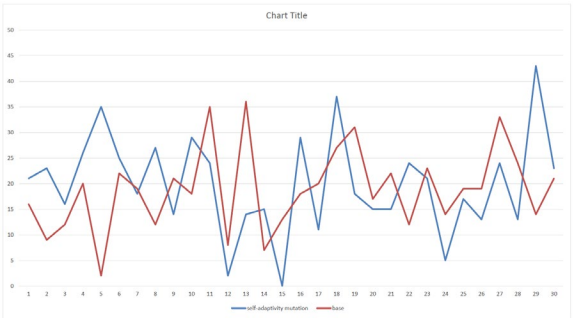
Variable 1	Variable 2
Mean	19.9
Variance	92.78275862
Observations	30
df	29
F	1.402282677
P(F<=f) one-tail	0.183944101
F Critical one-tail	1.860811435

F < F Critical

assume unequal variances

t-Test: Two-Sample Assuming Unequal Variances

Variable 1	Variable 2
Mean	19.9
Variance	92.78275862
Observations	30
Hypothesized Mean Difference	0
df	56
t Stat	0.477887204
P(T<=t) one-tail	0.31729535
t Critical one-tail	1.67252303
P(T<=t) two-tail	0.634590699
t Critical two-tail	2.003240719



Red 2: Self-adaptive Penalty Coefficient

For the self-adaptive penalty coefficient, it was essentially the same as the self-adaptive mutation rate. A new gene was added to allow the penalty coefficient to be defined and to be evolved per individual. I do not think a self-adaptive penalty coefficient sounds like a good idea, however. Since we no longer throw away invalid solutions, invalid solutions can thrive with a low penalty coefficient as they will look, as far as the fitness algorithm is concerned, as valid as “valid” solutions. Therefore, the most efficient penalty coefficient would be 0, allowing extremely invalid solutions to have high fitness values. Since our logging only records overall and average fitness one cannot see the actual board of these individuals but I would assume that most of them are producing extremely invalid solutions; such as intersecting rays or unsatisfied black cells.

C1

```
Base self-adaptivity penalty
35 98.51349233
29 98.76598274
12 99.93154783
17 96.43673276
25 98.89764611
10 97.27910252
21 99.68579623
12 97.62778949
21 98.88338121
20 99.14742977
20 97.40415555
7 98.50086958
25 99.23223584
35 98.42612806
29 98.9276992
20 98.02070913
39 95.90952239
42 98.6019246
43 99.57625481
19 98.16446438
29 99.17504075
33 98.28261915
8 97.69119086
28 99.95488727
11 99.77693392
7 99.3815715
36 96.94309693
29 98.33545797
28 97.80100989
33 98.9341456
```

F-Test Two-Sample for Variances

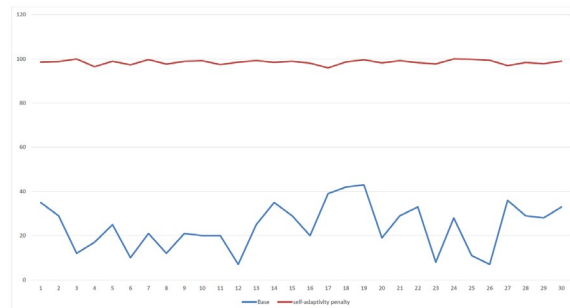
	Variable 1	Variable 2
Mean	24.1	98.47362728
Variance	110.9896552	1.023959678
Observations	30	30
df	29	29
F	108.3926033	
P(F<=f) one-tail	9.55367E-23	
F Critical one-tail	1.860811435	

F < F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	24.1	98.47362728
Variance	110.9896552	1.023959678
Observations	30	30
Pooled Variance	56.00680743	
Hypothesized Mean Difference	0	
df	58	
t Stat	-38.48966968	
P(T<=t) one-tail	2.69868E-43	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	5.39737E-43	
t Critical two-tail	2.001717484	



C2

```
base self-adaptivity penalty
16 98.58708902
9 97.77089032
12 99.86971713
20 99.39419861
2 97.51339485
22 99.40958748
19 98.89313062
12 99.84415584
21 99.73001833
18 99.97978017
35 98.81741526
8 99.6196266
36 99.12341946
7 98.86841303
13 98.09245905
18 99.29847408
20 97.07634327
27 97.3166693
31 98.81906314
17 99.92182287
22 97.71806091
12 98.64215736
23 99.19298544
14 97.83429786
19 99.77617159
19 95.28961942
33 98.5435956
24 98.22118479
14 99.38474406
21 97.69944911
```

F-Test Two-Sample for Variances

	Variable 1	Variable 2
Mean	18.8	98.67481267
Variance	66.16551724	1.137243853
Observations	30	30
df	29	29
F	58.18058904	
P(F<=f) one-tail	6.39529E-19	
F Critical one-tail	1.860811435	

F > F Critical

assume equal variance

t-Test: Two-Sample Assuming Equal Variances

	Variable 1	Variable 2
Mean	18.8	98.67481267
Variance	66.16551724	1.137243853
Observations	30	30
Pooled Variance	33.65138055	
Hypothesized Mean Difference	0	
df	58	
t Stat	-53.32784564	
P(T<=t) one-tail	2.77562E-51	
t Critical one-tail	1.671552762	
P(T<=t) two-tail	5.55123E-51	
t Critical two-tail	2.001717484	

